

Auto-Scaled Service

Objective

Launch an auto-scaled service in Amazon's Virtual Private Cloud.

Be creative! This is your chance to showcase your experience with a particular language or tool. Ideally, though, this should take no more than three days.

Guidelines

1. Create an executable script for deploying your service. This script should be configurable with command line parameters and should return the IP or DNS name of your service. At a minimum you'll need this script to create the following Amazon resources to support your service:
 - a. Autoscaling Group
 - b. Launch Configuration
 - c. Security Group(s)
 - d. Elastic Load Balancer (don't worry about ssl)
2. Automagically provision your instances with the tool / language of your choice. Each instance should serve a static page with nginx.
3. Ensure that your service has sensible security group(s) and is publicly available.
4. Ensure that your service can withstand the termination of a node and auto-heal.

Note: When your script is ran, it needs to successfully create and provision the necessary resources for your homework to be accepted.

Deliverables

GitHub repo containing your launch script, provisioning code and any other artifacts.

Extra Credit

Use CloudWatch metrics to automatically scale your service based on CPU or request load.